

COMPSCIX 415.2 Homework 9/Final

Rajat Jain

Aug 4th, 2018

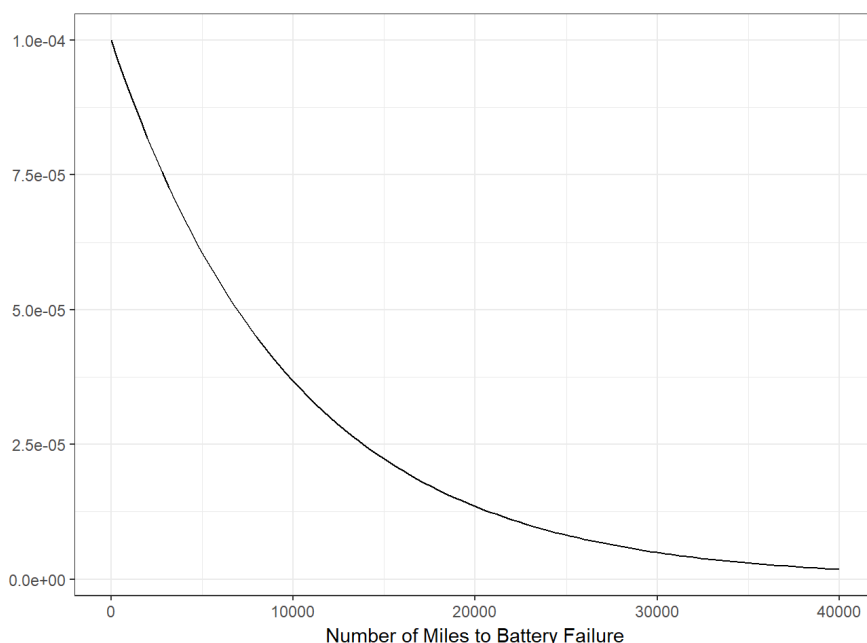
Contents

Exercise 1 - Sampling Distributions, Functions and For Loops (10 points)	1
Exercise 2 - Linear Regression (5 points)	6
Exercise 3 - AB Testing (5 points)	7

Exercise 1 - Sampling Distributions, Functions and For Loops (10 points)

Recall that the distribution of the sample mean is approximately a Normal distribution, and that the standard error is $\frac{\sigma}{\sqrt{n}}$. This holds true regardless of the distribution of our population.

For this problem, assume that the number of miles that a particular car can run before its battery wears out is exponentially distributed with an average of 10,000 miles. The exponential distribution looks like this:



The exponential distribution has a rate parameter that controls how quickly the distribution decays and defines what the mean and standard deviation will be. In our case the **rate** = $1/10000$, the **mean** = 10000 and the **standard deviation** = 10000. You can sample from this exponential distribution in R using this code:

```
# sample size
samp_size <- 100
# set the rate parameter
samp_rate <- 1/10000

# take sample
rexp(n = samp_size, rate = samp_rate)
```

STEP 1

Write an R function that does the following:

- Takes a sample of size `samp_size` from this exponential distribution (`samp_size` is an input parameter for the function)
- Calculates the mean of that sample
- Calculates the standard deviation of that sample
- Returns the calculated mean and standard deviation as a list

Helper code

```
samp_fun <- function(samp_size, samp_rate) {  
  
  ...your code here...  
  
  stats <- list(samp_avg = samp_avg, samp_std_dev = samp_std_dev)  
  return(stats)  
}
```

Here is the code for the desired function:

```
samp_stats <- function(samp_size, samp_rate) {  
  
  sample <- rexp(n = samp_size, rate = samp_rate)  
  
  samp_avg <- mean(sample)  
  samp_sd <- sd(sample)  
  
  stats <- list(samp_avg = samp_avg, samp_std_dev = samp_sd)  
  return(stats)  
}
```

STEP 2

Then write a loop that does this:

- Runs the above function 1000 times, with `samp_size = 50` and `samp_rate = 1/10000`
- Saves all of the sample means in a vector called `sample_means`, and all of the sample standard deviations in a vector called `sample_sds`


```
# No. of iterations.  
N <- 1000  
# Sample size.  
n <- 50  
  
# Output Vectors.  
sample_means <- rep(NA, N)  
sample_sds <- rep(NA, N)  
  
for (i in 1:N) {  
  stats <- samp_stats(n, 1/10000)  
  sample_means[i] <- stats$samp_avg  
  sample_sds[i] <- stats$samp_std_dev  
}
```

STEP 3

Then

- plot your sample means as a histogram

```
ggplot(data.frame(samp_avg = sample_means), aes(x = samp_avg)) +  
  geom_histogram(binwidth = 200) + xlab("Sample Averages") + theme_bw()
```



- output the standard deviation of your sample means

```
paste("Standard Deviation of sample means is", round(sd(sample_means), 2))
```

```
## [1] "Standard Deviation of sample means is 1398.41"
```

- calculate the theoretical standard error ($\sigma = 10000$, $n = \text{sample size}$)

```
paste("Theoretical Standard Error is", round(10000 / sqrt(n), 2))
```

```
## [1] "Theoretical Standard Error is 1414.21"
```

- calculate the mean of the sample standard deviations and use this to calculate the empirical standard error

```
paste("Mean of sample Standard Deviations is", round(mean(sample_sds), 2),  
      "and Emperical Standard Error is", round(mean(sample_sds) / sqrt(n), 2))
```

```
## [1] "Mean of sample Standard Deviations is 9849.74 and Emperical Standard Error is 1392.96"
```

STEP 4


Repeat STEP 2 and STEP 3 using a sample size of 500 and 5000.

Repeating STEP 2 and STEP 3 using a sample size of 500:

```
n <- 500 # Sample size.
sample_means <- sample_sds <- rep(NA, N) # Output Vectors

# Run the function 1000 times, with samp_size = 500 and samp_rate = 1/10000
for (i in 1:N) {
  stats <- samp_stats(n, 1/10000)
  sample_means[i] <- stats$samp_avg
  sample_sds[i] <- stats$samp_std_dev
}

# Plot the sample means as a histogram
ggplot(data.frame(samp_avg = sample_means), aes(x = samp_avg)) +
  geom_histogram(fill = "red") + xlab("Sample Averages") + theme_bw()
```



```
# Output the standard deviation of the sample means
paste("Standard Deviation of sample means is", round(sd(sample_means), 2))
```

```
## [1] "Standard Deviation of sample means is 443.41"
```

```
# Calculate the theoretical standard error
paste("Theoretical Standard Error is", round(10000 / sqrt(n), 2))
```

```
## [1] "Theoretical Standard Error is 447.21"
```

```
# Calculate the empirical standard error
paste("Mean of sample Standard Deviations is", round(mean(sample_sds), 2),
      "and Emperical Standard Error is", round(mean(sample_sds) / sqrt(n), 2))
```


```
## [1] "Mean of sample Standard Deviations is 9971.02 and Emperical Standard Error is 445.92"
```

Repeating STEP 2 and STEP 3 using a sample size of 5000:

```
n <- 5000 # Sample size.
sample_means <- sample_sds <- rep(NA, N) # Output Vectors

# Run the function 1000 times, with samp_size = 5000 and samp_rate = 1/10000
for (i in 1:N) {
  stats <- samp_stats(n, 1/10000)
  sample_means[i] <- stats$samp_avg
  sample_sds[i] <- stats$samp_std_dev
}

# Plot the sample means as a histogram
ggplot(data.frame(samp_avg = sample_means), aes(x = samp_avg)) +
  geom_histogram(fill = "green") + xlab("Sample Averages") + theme_bw()
```



```
# Output the standard deviation of the sample means
paste("Standard Deviation of sample means is", round(sd(sample_means), 2))
```

```
## [1] "Standard Deviation of sample means is 142.46"
```

```
# Calculate the theoretical standard error
paste("Theoretical Standard Error is", round(10000 / sqrt(n), 2))
```

```
## [1] "Theoretical Standard Error is 141.42"
```

```
# Calculate the empirical standard error
paste("Mean of sample Standard Deviations is", round(mean(sample_sds), 2),
      "and Emperical Standard Error is", round(mean(sample_sds) / sqrt(n), 2))
```

```
## [1] "Mean of sample Standard Deviations is 9997.92 and Emperical Standard Error is 141.39"
```

Exercise 2 - Linear Regression (5 points)

For this exercise we will return to the House Prices prediction dataset that we used for HW 7. You should have already downloaded the `train.csv` dataset before, but if you didn't you can download it from Canvas in this week's module.

Load the `train.csv` dataset into R and fit a regression model with:

- `y = SalePrice`
- Features: `LotArea`, `OverallQual`, and `ExterQual`

Here is the code to load data and fit a linear regression model:

```
train_data <- read_csv('train.csv')
model <- lm(formula = SalePrice ~ LotArea + OverallQual + ExterQual, data = train_data)
```

Answer these questions:

- Use the broom package to output the coefficients and the R-squared

```
tidy(model)
```

```
## # A tibble: 6 x 5
##   term          estimate std.error statistic    p.value
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)  40764.    12358.      3.30 9.95e- 4
## 2 LotArea       1.45      0.116     12.5 3.72e- 34
## 3 OverallQual  34466.     1216.     28.3 4.77e-141
## 4 ExterQualFa -95352.     14592.    -6.53 8.80e- 11
## 5 ExterQualGd -71529.      6737.    -10.6 2.05e- 25
## 6 ExterQualTA -97527.      7541.    -12.9 2.72e- 36
```

```
glance(model)
```

```
## # A tibble: 1 x 11
##   r.squared adj.r.squared sigma statistic p.value    df logLik    AIC
##   <dbl>      <dbl>    <dbl>    <dbl>    <dbl> <int>  <dbl>  <dbl>
## 1    0.695      0.694 43948.      663.      0.      6 -17677. 35368.
## # ... with 3 more variables: BIC <dbl>, deviance <dbl>, df.residual <int>
```

- Interpret the coefficient on `LotArea`

Controlling for all other features, if the `LotArea` increases by 1 unit, the `SalePrice` increases by \$1.45 on average.

- Interpret the coefficient on `ExterQualGd`

`SalePrice` for houses with `ExterQualGd` (Good exterior material quality) are on average \$71,529.49 lower relative to the houses with `ExterQualEx` (Excellent) while controlling for all other features. *PS: There are no entries in training data with `ExterQual` value `Po` (Poor).*

- Compare this model to the model we fit in HW 7 with `GrLivArea`, `OverallQual`, `Neighborhood`. Which is the better fitting model?

This model has the adjusted-R-squared value of 0.69 which is lower than the adjusted-R-squared value of 0.78 achieved with the model we fit in HW 7 with `GrLivArea`, `OverallQual`, `Neighborhood`. It means that the current model explains less variability in the data as compared to the one built in HW 7. Hence, the model from HW 7 was a better fitting model.

Exercise 3 - AB Testing (5 points)

Download the `ab_test_data.csv` file from Canvas. This file contains two columns: `version` and `conversion`. Each row is a visitor to a webpage. The `version` column tells us which version of the webpage the visitor saw, and the `conversion` column is a binary value and equals 1 if the visitor converted (0 otherwise).

We want to perform an AB test on this data to see if the conversion rates are different for the two versions of the webpage.

Reading the data file:

```
ab_test_data <- read_csv('ab_test_data.csv')
```

```
## Parsed with column specification:
## cols(
##   version = col_character(),
##   conversion = col_integer()
## )
```

Answer these questions:

a. What proportion of visitors converted for each version of the webpage?

```
ab_test_data %>%
  group_by(version) %>%
  summarise(
    conversion_rate = 100 * mean(conversion),
    num_converted   = sum(conversion),
    num_visited     = n()
  )
```

```
## # A tibble: 2 x 4
##   version conversion_rate num_converted num_visited
##   <chr>         <dbl>         <int>         <int>
## 1 A             4.15             83           2000
## 2 B            10.0            200           2000
```

4.15% visitors converted on version A and 10% visitors converted on version B of the webpage.

b. Perform the AB test in R. What is the p-value for the AB test (hypothesis test of proportions)?

Performing the A/B test using proportion test:

```
test_results <- prop.test(c(83, 200), c(2000, 2000))
```

```
# p-value
test_results$p.value
```

```
## [1] 8.479709e-13
```

Since the p-value for the A/B test is less than 0.05 (α), we can reject the Null Hypothesis that version A and B of the webpage have same conversion rate. This means that our alternate hypothesis is accepted and version A and B have statistically significant differences in conversion rates.