

COMPSCIX 415.2 Homework 3

Rajat Jain

June 25th, 2018


Contents

Section 3.8.1 Exercises	1
Section 3.9.1 Exercises	4
Section 4.4 Exercises	5
Section 5.2.4 Exercises	7
Section 5.4.1 Exercises	12

Section 3.8.1 Exercises


1. What is the problem with this plot? How could you improve it?

```
ggplot(data = mpg, mapping = aes(x = cty, y = hwy)) +  
  geom_point()
```



This plot has many overlapping point which do not show the areas of concentration properly. It can be improved by adding jitter to the position adjustment.

```
ggplot(data = mpg, mapping = aes(x = cty, y = hwy)) +  
  geom_point(position = "jitter")
```



2. What parameters to `geom_jitter()` control the amount of jittering?

The parameters to `geom_jitter()` which control the amount of jittering are `width` and `height`.

3. Compare and contrast `geom_jitter()` with `geom_count()`.

`geom_jitter` randomly moves the overlapping points slightly to avoid overlapping whereas, `geom_count` counts the overlapping points at a given point and maps them to the size of a single point. The determinism of `geom_count` makes it useful in discrete situations, but it does not work when the points are not exactly overlapping but are very close.

4. What's the default position adjustment for `geom_boxplot()`? Create a visualisation of the mpg dataset that demonstrates it.

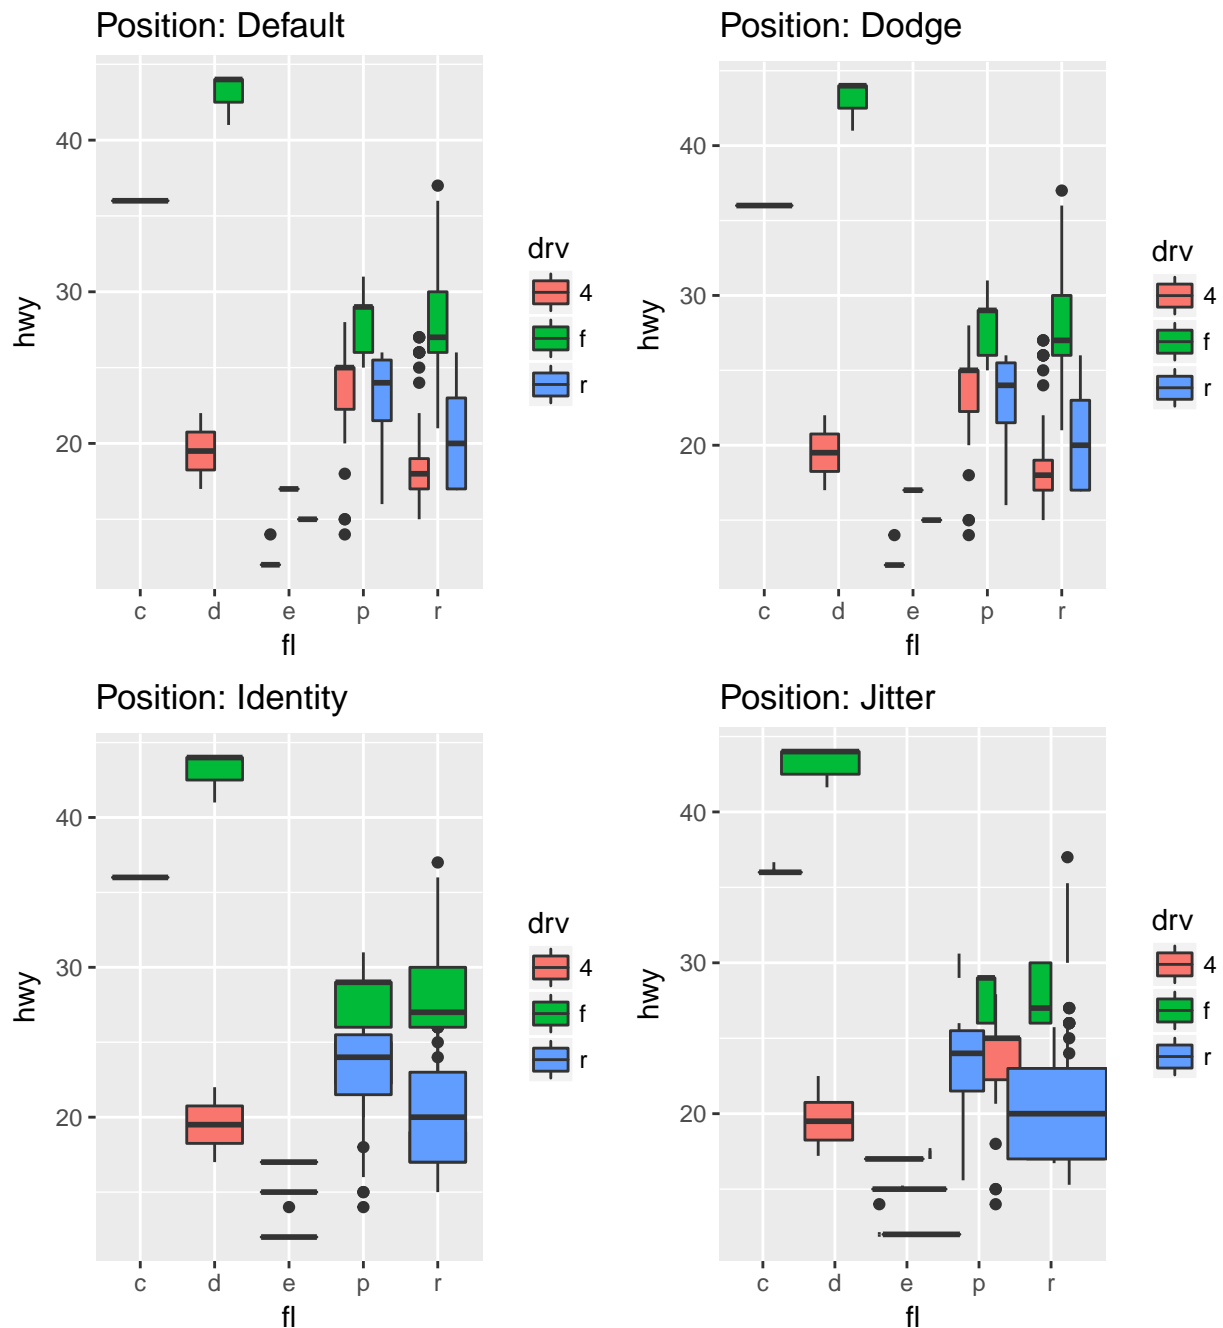
The default position adjustment for `geom_boxplot()` is `position_dodge`. Here is a demonstration:

```
my_plot <- ggplot(data = mpg, mapping = aes(x = fl, y = hwy, fill = drv))  
  
# Plot with default position adjustment.  
dflt <- my_plot + ggtitle("Position: Default") +  
  geom_boxplot()  
  
# Plot with dodge position adjustment.  
ddge <- my_plot + ggtitle("Position: Dodge") +  
  geom_boxplot(position = "dodge")
```

```
# Plot with identity position adjustment.
idnt <- my_plot + ggtitle("Position: Identity") +
  geom_boxplot(position = "identity")

# Plot with jitter position adjustment.
jter <- my_plot + ggtitle("Position: Jitter") +
  geom_boxplot(position = "jitter")

grid.arrange(dflt, ddge, idnt, jter, nrow = 2, ncol = 2)
```



In the above plots position adjustment **Dodge** generated the plot looking exactly as with default position adjustment.


Section 3.9.1 Exercises

2. What does `labs()` do? Read the documentation.

`labs()` function from `ggplot2` package is used to modify axis, legend, and plot labels.

4. What does the plot below tell you about the relationship between city and highway mpg? Why is `coord_fixed()` important? What does `geom_abline()` do?

```
ggplot(data = mpg, mapping = aes(x = cty, y = hwy)) +  
  geom_point() +  
  geom_abline() +  
  coord_fixed()
```



This plot tells us that the city and highway mpg are positively correlated, meaning if a vehicle has a higher city mpg, it is also expected to have a higher highway mpg.

`coord_fixed()` is important because both city and highway mpg have the same unit and hence one unit of each should be represented by the same length on the axes for the best visualization. `coord_fixed()` does exactly that for us.

`geom_abline()` creates a reference line shown on the plot. Since it is called without `intercept` and `slope` arguments, it generates a default reference line passing through the origin and with a slope of 1 (45°). So this line represents all the point on the plot where city mpg would be equal to highway mpg. From the plot we can also observe that all the plotted points are above this line, which means that the highway mpg is always higher than the city mpg. If we observe closely, we can also see that the scatter plot is more or less parallel to the reference line. This tells us that the highway mpg is higher than the city mpg for a constant additive offset.

Section 4.4 Exercises

1. Why does this code not work?

```
my_variable <- 10
my_variable
#> Error in eval(expr, envir, enclos): object 'my_variable' not found``
```

Look carefully! (This may seem like an exercise in pointlessness, but training your brain to notice even the tiniest difference will pay off when programming.)

This code does not work because the name of the variable while printing it is mis-spelled.

```
my_variable <- 10
my_variable
```

```
## [1] 10
```

The above code is its fixed version which works.

2. Tweak each of the following R commands so that they run correctly:


```
library(tidyverse)

ggplot(dota = mpg) +
  geom_point(mapping = aes(x = displ, y = hwy))

fliter(mpg, cyl = 8)
filter(diamond, carat > 3)
```

Here is the fixed and working code:

```
# Package tidyverse is pre-loaded in this document.
ggplot(data = mpg) +
  geom_point(mapping = aes(x = displ, y = hwy))
```



```
filter(mpg, cyl == 8)
```

```
## # A tibble: 70 x 11
##   manufacturer model      displ  year  cyl trans  drv      cty   hwy fl
##   <chr>          <chr>    <dbl> <int> <int> <chr> <chr> <int> <int> <chr>
## 1 audi          a6 quatt~ 4.20  2008    8 auto(~ 4      16    23 p
## 2 chevrolet     c1500 su~ 5.30  2008    8 auto(~ r      14    20 r
## 3 chevrolet     c1500 su~ 5.30  2008    8 auto(~ r      11    15 e
## 4 chevrolet     c1500 su~ 5.30  2008    8 auto(~ r      14    20 r
## 5 chevrolet     c1500 su~ 5.70  1999    8 auto(~ r      13    17 r
## 6 chevrolet     c1500 su~ 6.00  2008    8 auto(~ r      12    17 r
## 7 chevrolet     corvette  5.70  1999    8 manua~ r      16    26 p
## 8 chevrolet     corvette  5.70  1999    8 auto(~ r      15    23 p
## 9 chevrolet     corvette  6.20  2008    8 manua~ r      16    26 p
## 10 chevrolet    corvette  6.20  2008    8 auto(~ r      15    25 p
## # ... with 60 more rows, and 1 more variable: class <chr>
```

```
filter(diamonds, carat > 3)
```

```
## # A tibble: 32 x 10
##   carat cut      color clarity depth table price      x      y      z
##   <dbl> <ord>    <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl> <dbl>
## 1 3.01 Premium I      I1      62.7  58.  8040  9.10  8.97  5.67
## 2 3.11 Fair J        I1      65.9  57.  9823  9.15  9.02  5.98
## 3 3.01 Premium F      I1      62.2  56.  9925  9.24  9.13  5.73
## 4 3.05 Premium E      I1      60.9  58. 10453  9.26  9.25  5.66
## 5 3.02 Fair I        I1      65.2  56. 10577  9.11  9.02  5.91
## 6 3.01 Fair H        I1      56.1  62. 10761  9.54  9.38  5.31
## 7 3.65 Fair H        I1      67.1  53. 11668  9.53  9.48  6.38
## 8 3.24 Premium H      I1      62.1  58. 12300  9.44  9.40  5.85
## 9 3.22 Ideal I        I1      62.6  55. 12545  9.49  9.42  5.92
## 10 3.50 Ideal H        I1      62.8  57. 12587  9.65  9.59  6.03
## # ... with 22 more rows
```

Section 5.2.4 Exercises

1. Find all flights that

1. Had an arrival delay of two or more hours

```
filter(flights, arr_delay >= 120)
```

```
## # A tibble: 10,200 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     811             630          101.    1047
## 2  2013     1     1     848             1835         853.    1001
## 3  2013     1     1     957             733          144.    1056
## 4  2013     1     1    1114             900          134.    1447
## 5  2013     1     1    1505             1310         115.    1638
## 6  2013     1     1    1525             1340         105.    1831
## 7  2013     1     1    1549             1445           64.    1912
## 8  2013     1     1    1558             1359         119.    1718
## 9  2013     1     1    1732             1630           62.    2028
## 10 2013     1     1    1803             1620         103.    2008
## # ... with 10,190 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

2. Flew to Houston (IAH or HOU)

```
filter(flights, dest %in% c("IAH", "HOU"))
```

```
## # A tibble: 9,313 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517             515           2.     830
## 2  2013     1     1     533             529           4.     850
## 3  2013     1     1     623             627          -4.     933
## 4  2013     1     1     728             732          -4.    1041
## 5  2013     1     1     739             739           0.    1104
## 6  2013     1     1     908             908           0.    1228
## 7  2013     1     1    1028            1026           2.    1350
## 8  2013     1     1    1044            1045          -1.    1352
## 9  2013     1     1    1114             900         134.    1447
## 10 2013     1     1    1205            1200           5.    1503
## # ... with 9,303 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

3. Were operated by United, American, or Delta

```
filter(flights, carrier %in% c("UA", "AA", "DL"))
```

```
## # A tibble: 139,504 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
```

```
## 1 2013 1 1 517 515 2. 830
## 2 2013 1 1 533 529 4. 850
## 3 2013 1 1 542 540 2. 923
## 4 2013 1 1 554 600 -6. 812
## 5 2013 1 1 554 558 -4. 740
## 6 2013 1 1 558 600 -2. 753
## 7 2013 1 1 558 600 -2. 924
## 8 2013 1 1 558 600 -2. 923
## 9 2013 1 1 559 600 -1. 941
## 10 2013 1 1 559 600 -1. 854
## # ... with 139,494 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

4. Departed in summer (July, August, and September)

```
filter(flights, month %in% 7:9)
```

```
## # A tibble: 86,326 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1 2013     7     1       1         2029         212.    236
## 2 2013     7     1       2         2359          3.    344
## 3 2013     7     1      29         2245        104.    151
## 4 2013     7     1      43         2130        193.    322
## 5 2013     7     1      44         2150        174.    300
## 6 2013     7     1      46         2051        235.    304
## 7 2013     7     1      48         2001        287.    308
## 8 2013     7     1      58         2155        183.    335
## 9 2013     7     1     100         2146        194.    327
## 10 2013     7     1     100         2245        135.    337
## # ... with 86,316 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

5. Arrived more than two hours late, but didn't leave late

```
filter(flights, arr_delay > 120 & dep_delay <= 0)
```

```
## # A tibble: 29 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1 2013     1    27    1419         1420         -1.    1754
## 2 2013    10     7    1350         1350          0.    1736
## 3 2013    10     7    1357         1359         -2.    1858
## 4 2013    10    16     657          700         -3.    1258
## 5 2013    11     1     658          700         -2.    1329
## 6 2013     3    18    1844         1847         -3.     39
## 7 2013     4    17    1635         1640         -5.    2049
## 8 2013     4    18     558          600         -2.    1149
## 9 2013     4    18     655          700         -5.    1213
```



```
## 10 2013      5      22      1827      1830      -3.      2217
## # ... with 19 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

6. Were delayed by at least an hour, but made up over 30 minutes in flight

```
filter(flights, dep_delay >= 60 & dep_delay - arr_delay > 30)
```

```
## # A tibble: 1,844 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1    2205           1720         285.     46
## 2  2013     1     1    2326           2130         116.    131
## 3  2013     1     3    1503           1221         162.   1803
## 4  2013     1     3    1839           1700          99.   2056
## 5  2013     1     3    1850           1745          65.   2148
## 6  2013     1     3    1941           1759         102.   2246
## 7  2013     1     3    1950           1845          65.   2228
## 8  2013     1     3    2015           1915          60.   2135
## 9  2013     1     3    2257           2000         177.     45
## 10 2013     1     4    1917           1700         137.   2135
## # ... with 1,834 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

7. Departed between midnight and 6am (inclusive)

```
filter(flights, dep_time <= 600 | dep_time == 2400)
```

```
## # A tibble: 9,373 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517           515           2.     830
## 2  2013     1     1     533           529           4.     850
## 3  2013     1     1     542           540           2.     923
## 4  2013     1     1     544           545          -1.   1004
## 5  2013     1     1     554           600          -6.     812
## 6  2013     1     1     554           558          -4.     740
## 7  2013     1     1     555           600          -5.     913
## 8  2013     1     1     557           600          -3.     709
## 9  2013     1     1     557           600          -3.     838
## 10 2013     1     1     558           600          -2.     753
## # ... with 9,363 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

3. How many flights have a missing dep_time? What other variables are missing? What might these rows represent?

```
summary(flights)
```

```
##      year      month      day      dep_time
## Min.   :2013   Min.   : 1.000   Min.   : 1.00   Min.   :    1
## 1st Qu.:2013   1st Qu.: 4.000   1st Qu.: 8.00   1st Qu.: 907
## Median :2013   Median : 7.000   Median :16.00   Median :1401
## Mean   :2013   Mean   : 6.549   Mean   :15.71   Mean   :1349
## 3rd Qu.:2013   3rd Qu.:10.000   3rd Qu.:23.00   3rd Qu.:1744
## Max.   :2013   Max.   :12.000   Max.   :31.00   Max.   :2400
##                                     NA's   :8255
## sched_dep_time  dep_delay      arr_time  sched_arr_time
## Min.   : 106   Min.   : -43.00   Min.   :    1   Min.   :    1
## 1st Qu.: 906   1st Qu.:  -5.00   1st Qu.:1104   1st Qu.:1124
## Median :1359   Median :  -2.00   Median :1535   Median :1556
## Mean   :1344   Mean   : 12.64   Mean   :1502   Mean   :1536
## 3rd Qu.:1729   3rd Qu.: 11.00   3rd Qu.:1940   3rd Qu.:1945
## Max.   :2359   Max.   :1301.00   Max.   :2400   Max.   :2359
##                                     NA's   :8255      NA's   :8713
##      arr_delay      carrier      flight      tailnum
## Min.   : -86.000   Length:336776   Min.   :    1   Length:336776
## 1st Qu.: -17.000   Class :character 1st Qu.: 553   Class :character
## Median :  -5.000   Mode  :character Median :1496   Mode  :character
## Mean    :   6.895                                     Mean   :1972
## 3rd Qu.:  14.000                                     3rd Qu.:3465
## Max.    :1272.000                                     Max.   :8500
## NA's    :9430
##      origin      dest      air_time      distance
## Length:336776   Length:336776   Min.   : 20.0   Min.   : 17
## Class :character Class :character 1st Qu.: 82.0   1st Qu.: 502
## Mode  :character Mode  :character Median :129.0   Median : 872
##                                     Mean   :150.7   Mean   :1040
##                                     3rd Qu.:192.0   3rd Qu.:1389
##                                     Max.   :695.0   Max.   :4983
##                                     NA's   :9430
##      hour      minute      time_hour
## Min.   : 1.00   Min.   : 0.00   Min.   :2013-01-01 05:00:00
## 1st Qu.: 9.00   1st Qu.: 8.00   1st Qu.:2013-04-04 13:00:00
## Median :13.00   Median :29.00   Median :2013-07-03 10:00:00
## Mean   :13.18   Mean   :26.23   Mean   :2013-07-03 05:02:36
## 3rd Qu.:17.00   3rd Qu.:44.00   3rd Qu.:2013-10-01 07:00:00
## Max.   :23.00   Max.   :59.00   Max.   :2013-12-31 23:00:00
##
```

8255 flights have missing dep_time. Other variables with missing values are: dep_delay, arr_time, arr_delay & air_time.

Since only those variables seem to have missing values which are related to actual flight instances and not just the scheduled details, they may represent the cancelled flights. Alternatively, they could just be errors in data entry while recording those values at flight departure and/or arrival. Or a mix of both.

4. Why is $\text{NA} \wedge 0$ not missing? Why is $\text{NA} \mid \text{TRUE}$ not missing? Why is $\text{FALSE} \& \text{NA}$ not missing? Can you figure out the general rule? ($\text{NA} * 0$ is a tricky counterexample!)

```
NA ^ 0
```

```
## [1] 1
```

```
NA | TRUE
```

```
## [1] TRUE
```

```
FALSE & NA
```

```
## [1] FALSE
```

```
NA * 0
```

```
## [1] NA
```

$\text{NA} \wedge 0$ is not missing because anything to the power of 0 is 1 irrespective of what the value is. However, this is not true for $\text{Inf} \wedge 0$. It is still indeterminate form. So this seems to be an incorrect evaluation by R.

$\text{NA} \mid \text{TRUE}$ is not missing because the \mid expression evaluates to **TRUE** if either parts of the expression evaluate to **TRUE** and since **TRUE** is always **TRUE** it doesn't matter what the other part is.

Similarly, $\text{FALSE} \& \text{NA}$ is not missing because $\&$ expression evaluates to **FALSE** if either parts of the expression evaluates to **FALSE** irrespective of the value of the other part.

$\text{NA} * 0$ evaluates to **NA** because **NA** could take any value including **Inf** and $\text{Inf} * 0$ is indeterminate form.

Section 5.4.1 Exercises

1. Brainstorm as many ways as possible to select `dep_time`, `dep_delay`, `arr_time`, and `arr_delay` from `flights`.

```
# Define column names vector.
```

```
col_names <- c("dep_time", "dep_delay", "arr_time", "arr_delay")
```

```
#1. Using [] with column indexes.
```

```
flights[, c(4, 6, 7, 9)]
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>    <dbl>    <int>    <dbl>
## 1     517         2.     830        11.
## 2     533         4.     850        20.
## 3     542         2.     923        33.
## 4     544        -1.    1004       -18.
## 5     554        -6.     812       -25.
## 6     554        -4.     740        12.
## 7     555        -5.     913        19.
## 8     557        -3.     709       -14.
## 9     557        -3.     838         -8.
## 10    558        -2.     753         8.
## # ... with 336,766 more rows
```

```
#2. Using [] with column names vector.
```

```
flights[, col_names]
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>    <dbl>    <int>    <dbl>
## 1     517         2.     830        11.
## 2     533         4.     850        20.
## 3     542         2.     923        33.
## 4     544        -1.    1004       -18.
## 5     554        -6.     812       -25.
## 6     554        -4.     740        12.
## 7     555        -5.     913        19.
## 8     557        -3.     709       -14.
## 9     557        -3.     838         -8.
## 10    558        -2.     753         8.
## # ... with 336,766 more rows
```

```
#3. Using select with column indexes.
```

```
select(flights, 4, 6, 7, 9)
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>    <dbl>    <int>    <dbl>
## 1     517         2.     830        11.
## 2     533         4.     850        20.
## 3     542         2.     923        33.
## 4     544        -1.    1004       -18.
## 5     554        -6.     812       -25.
```

```
## 6      554      -4.      740      12.
## 7      555      -5.      913      19.
## 8      557      -3.      709     -14.
## 9      557      -3.      838      -8.
## 10     558      -2.      753       8.
## # ... with 336,766 more rows
```

```
#4. Using select with column names vector.
select(flights, col_names)
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>   <int>   <dbl>
## 1     517         2.     830     11.
## 2     533         4.     850     20.
## 3     542         2.     923     33.
## 4     544        -1.    1004    -18.
## 5     554        -6.     812    -25.
## 6     554        -4.     740     12.
## 7     555        -5.     913     19.
## 8     557        -3.     709    -14.
## 9     557        -3.     838     -8.
## 10    558        -2.     753       8.
## # ... with 336,766 more rows
```

```
#5. Using select with column names mentioned as ... arguments.
select(flights, dep_time, dep_delay, arr_time, arr_delay)
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>   <int>   <dbl>
## 1     517         2.     830     11.
## 2     533         4.     850     20.
## 3     542         2.     923     33.
## 4     544        -1.    1004    -18.
## 5     554        -6.     812    -25.
## 6     554        -4.     740     12.
## 7     555        -5.     913     19.
## 8     557        -3.     709    -14.
## 9     557        -3.     838     -8.
## 10    558        -2.     753       8.
## # ... with 336,766 more rows
```

```
#6. Using select with pipe operator and column indexes.
flights %>% select(4, 6, 7, 9)
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>   <int>   <dbl>
## 1     517         2.     830     11.
## 2     533         4.     850     20.
## 3     542         2.     923     33.
## 4     544        -1.    1004    -18.
## 5     554        -6.     812    -25.
## 6     554        -4.     740     12.
## 7     555        -5.     913     19.
```

```
## 8      557      -3.      709      -14.
## 9      557      -3.      838      -8.
## 10     558      -2.      753       8.
## # ... with 336,766 more rows
```

#7. Using select with pipe operator and column names vector.
 flights %>% select(col_names)

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>   <int>   <dbl>
## 1      517         2.     830     11.
## 2      533         4.     850     20.
## 3      542         2.     923     33.
## 4      544        -1.    1004    -18.
## 5      554        -6.     812    -25.
## 6      554        -4.     740     12.
## 7      555        -5.     913     19.
## 8      557        -3.     709    -14.
## 9      557        -3.     838     -8.
## 10     558        -2.     753      8.
## # ... with 336,766 more rows
```

#8. Using select with pipe operator and column names mentioned as ... arguments
 flights %>% select(dep_time, dep_delay, arr_time, arr_delay)

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>   <int>   <dbl>
## 1      517         2.     830     11.
## 2      533         4.     850     20.
## 3      542         2.     923     33.
## 4      544        -1.    1004    -18.
## 5      554        -6.     812    -25.
## 6      554        -4.     740     12.
## 7      555        -5.     913     19.
## 8      557        -3.     709    -14.
## 9      557        -3.     838     -8.
## 10     558        -2.     753      8.
## # ... with 336,766 more rows
```

#9. Using starts_with() select helper.
 flights %>% select(starts_with("dep_"), starts_with("arr_"))

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>   <int>   <dbl>
## 1      517         2.     830     11.
## 2      533         4.     850     20.
## 3      542         2.     923     33.
## 4      544        -1.    1004    -18.
## 5      554        -6.     812    -25.
## 6      554        -4.     740     12.
## 7      555        -5.     913     19.
## 8      557        -3.     709    -14.
## 9      557        -3.     838     -8.
```

```
## 10      558      -2.      753      8.
## # ... with 336,766 more rows
```

```
#10. Using one_of() select helper.
flights %>% select(one_of(col_names))
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>    <int>     <dbl>
## 1     517         2.     830         11.
## 2     533         4.     850         20.
## 3     542         2.     923         33.
## 4     544        -1.    1004        -18.
## 5     554        -6.     812        -25.
## 6     554        -4.     740         12.
## 7     555        -5.     913         19.
## 8     557        -3.     709        -14.
## 9     557        -3.     838         -8.
## 10     558        -2.     753          8.
## # ... with 336,766 more rows
```

3. What does the `one_of()` function do? Why might it be helpful in conjunction with this vector?

```
vars <- c("year", "month", "day", "dep_delay", "arr_delay")
```

`one_of()` is a Select Helper function from `dplyr` package. It allows for guessing or subset-matching. (Courtesy: StackOverflow). It returns all the columns with names which match the vector provided to it. For example, in conjunction with the vector above, it could be used with `select()` as below:

```
flights %>% select(one_of(vars))
```

```
## # A tibble: 336,776 x 5
##   year month   day dep_delay arr_delay
##   <int> <int> <int>     <dbl>     <dbl>
## 1  2013     1     1         2.         11.
## 2  2013     1     1         4.         20.
## 3  2013     1     1         2.         33.
## 4  2013     1     1        -1.        -18.
## 5  2013     1     1        -6.        -25.
## 6  2013     1     1        -4.         12.
## 7  2013     1     1        -5.         19.
## 8  2013     1     1        -3.        -14.
## 9  2013     1     1        -3.         -8.
## 10 2013     1     1        -2.          8.
## # ... with 336,766 more rows
```